

CS 4530: Fundamentals of Software Engineering

Module 4: Architecture of a Web Application

Adeel Bhutta and Mitch Wand
Khoury College of Computer Sciences


Learning Goals for this Lesson

At the end of this lesson, you should be able to


Web Applications are Distributed Systems

- Distributed systems are applications that perform computation across different computers
- Web applications are distributed systems *because*
 1. Users are literally in different places: you don't live in a data center
 2. Netflix won't fit entirely on one computer
- Distributed systems are hard!
 - Adopting well-understood patterns makes the task easier
 - We'll talk about one kind of pattern in this lecture
 - If you're Google, Netflix, or Amazon, the simple designs in this lecture won't serve your purposes

The simplest distributed system has a server and one or more clients



The server and the client communicate using an agreed-upon protocol




Examples of protocols

- Checking guests off a guest list
- Morse code
- Procedure calls
 - how are arguments passed to procedures?
 - how are results passed back to the caller?
- The system designer gets to choose the protocol, and then builds the programs to match


WebSockets is a useful protocol for event-driven applications

- A server can communicate with many clients
- A chat message is an “event” that the server wants everyone to know about.
- The web server is a *centralized* point that allows distributed clients to talk to each other.
 - Centralization is a cheat code for making distributed systems manageable



HTTP is a useful protocol for "function-call"-like patterns


HTTP gathers messages in pairs: client sends, server responds



HTTP is a protocol useful for “Function-call-like” Design Patterns


HTTP supports common design patterns:

- *Remote Procedure Call (RPC)*:
one computer is calling a function on another computer
- *REpresentational State Transfer (REST)*:
strictly client/server, server doesn't maintain a session or remember specific clients
- A set of routes (like /cities, /states, ...) answering HTTP requests is a “**REST API**”



http requests typically contain a **route** and a **payload**.

Here's a very simple express server




The Controller-Service-Repository Design

- This is the design we will use for all our web servers.
- It's only one way to organize a web server.
- But it's good enough for us.
- It emphasizes an important SE principle: Information Hiding
- If you are Google, you probably want something more sophisticated.

Three parts of our web servers

- The Controller-Service-Repository design separates the server into three parts
 - The **controller** knows how we communicate with client, doesn't know how we store data
 - The **service** doesn't know how we connect to the client or store data, but does as much of the interesting work ("business logic") as possible without knowing these things
 - The **repository** is the only part that stores information long-term



4.1 The Controller

- The controller is responsible for communication with the outside world
- We'll start by looking at the very simple controller we saw a few slides ago.
- Then we will add a little structure.

First, we add validation to the controller (here using Zod)

```
import express from "express";
import { type Request, type Response } from "express";
import { z } from "zod"
const app = express();
let numLogins = 0;
```

src/loginServer/withZod.ts

```
const zLogin = z.object({
  username: z.string().min(4),
  password: z.string().min(8)
})
```

Description of the
valid payloads



```
app.post('/api/user/login', (request:Request, response:Response) => {
  const validatedRequest = zLogin.safeParse(request)
  if (validatedRequest.error) {response.send({error: "Bad Request"})}
  else {
    const username:string = validatedRequest.data.username
    const password:string = validatedRequest.data.password
    // action goes here
  }
})
```

safeParse.data
returns correctly
typed components



The controller relies on the service layer to perform the appropriate actions

```
import express from "express";
import { type Request, type Response } from "express";
import { z } from "zod"
const app = express();
import { isAuthenticated, incrementLogins } from "../loginService.ts"

// validation goes here

const username:string = validatedRequest.data.username
const password:string = validatedRequest.data.password
if (isAuthenticated(username, password)) {
  const numLogins = incrementLogins()
  response.send({success:true, numLogins})
} else
  response.send({error: "Invalid username or password"})
}
```

The controller uses the service layer to assemble the data for an appropriate response

The controller sends the actual response itself

src/loginServer/withZodAndService.ts

4.2 The Service Layer Provides the Business Logic

```
var numLogins = 0

export function isAuthorized(username:string, password:string):boolean {
    return ((username == "user1") && (password == "secret"))
}

export function incrementLogins(): number {
    numLogins++;
    return numLogins
}
```

src/loginServer/loginService.ts


Looking back at the Transcript Server from the TDD Lecture

- Most of we saw from the transcript server was the **business logic** — the boring-sounding name that refers to most of the interesting stuff a web server does

But there was no clear separation of a repository layer

src/transcript.service.ts
(from M02 examples)

```
/** Adds a new student to the database
 * @param {string} newName - the name of the student
 * @returns {StudentID} - the newly-assigned ID for the new student
 */
addStudent(newName: string): StudentID {
  const newID = this._lastID++;
  const newStudent: Student = { studentID: newID,
                                studentName: newName };
  this._transcripts.push({
    student: newStudent, grades: []
  });
  return newID;
}
```



NOT business logic!

4.3 The Repository Layer

A Repository Layer should provide persistence

- Logins should be cumulative even if we restart the server
- Adding users and changing passwords shouldn't necessarily require updating code
- Lots of ways to achieve this:
 - MongoDB
 - PostgreSQL
 - SQLite
 - A file on the hard drive

Adding a persistent repository requires some new programming ideas

- Adding a persistent repository makes one big difference!
 - almost every action that reads or writes data is now *hundreds* of times slower, and involves reading to disk
 - this involves a relatively long delay, during which the CPU isn't doing anything useful
- JavaScript handles this with *asynchronous programming*.

Typescript uses the type system to distinguish long-running computations

- A long-running computation gets type **Promise<Grade>**, rather than **Grade**.

This is documented in the interface

```
export interface ITranscriptService {  
  addStudent(studentName: string): Promise<StudentID>;  
  getTranscript(id: StudentID): Promise<Transcript>;  
    // throws if id invalid  
  deleteStudent(id: StudentID): Promise<void>;  
    // throws if id invalid  
  addGrade(id: StudentID, course: Course, courseGrade: CourseGrade):  
    Promise<void>;  
  nameToIDs(studentName: string): Promise<StudentID[]>;  
}
```

src/withPersistence/ITranscriptService.ts

Use **await** to get the result of a promise (1)

```
describe('addStudent', () => {  
  it('should add a student to the database and return their id',  
    async () => {  
      expect(await service.nameToIDs('blair')).toStrictEqual([]);  
      const id1 = await service.addStudent('blair');  
      expect(await service.nameToIDs('blair')).toStrictEqual([id1]);  
    });  
});
```

- service.nameToIDs returns only a promise to compute a list of IDs
- here we need to see the actual list of IDs, so we need to wait for that computation to complete.

src/withPersistence/persistentService.spec.ts

Use **await** to get the result of a promise (2)

```
describe('addStudent', () => {  
  it('should add a student to the database and return their id',  
    async () => {  
      expect(await service.nameToIDs('blair')).toStrictEqual([]);  
      const id1 = await service.addStudent('blair');  
      expect(await service.nameToIDs('blair')).toStrictEqual([id1]);  
    });  
});
```

- So long as we have a single-threaded app, this works just fine!
- If our app is juggling multiple threads, there are dangers--more in M08.

Any function that calls a long-running function is itself long-running

```
async function numberOfSharedNames(name: string): Promise<number>
{
    const ids = await db.nameToIDs(name);
    return ids.length
}
```

- Of course, this means that any function that waits for the result of **numberOfSharedNames** must now also be async.
- Eventually, you will run out of functions-that-wait-for-an-async, so you will eventually run out of functions that need to be made async 😊

Review

At the end of this lesson, you should be able to

- Explain the role of “client” and “server” in the context of web application programming
- Describe the role of network protocols like http and websockets in a distributed system.
- Describe the fundamental differences between the three layers of the controller, service, and repository layers in a C-S-R architecture
- Use async and await to handle long-running procedures in a persistent repository